

- [PS1] Sasu Mäkinen, Henrik Skogström, Eero Laaksonen, and Tommi Mikkonen. Who needs mlops: What data scientists seek to accomplish and how can mlops help? In *2021 IEEE/ACM 1st Workshop on AI Engineering-Software Engineering for AI (WAIN)*, pages 109–112. IEEE, 2021.
- [PS2] Damian A Tamburri. Sustainable mlops: Trends and challenges. In *2020 22nd International Symposium on Symbolic and Numeric Algorithms for Scientific Computing (SYNASC)*, pages 17–23. IEEE, 2020.
- [PS3] Yue Zhou, Yue Yu, and Bo Ding. Towards mlops: A case study of ml pipeline platform. In *2020 International Conference on Artificial Intelligence and Computer Engineering (ICAICE)*, pages 494–500. IEEE, 2020.
- [PS4] Aquilas Tchanjou Njomou, Alexandra Johanne Bifona Africa, Bram Adams, and Marios Fokaefs. Msr4ml: Reconstructing artifact traceability in machine learning repositories. In *2021 IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER)*, pages 536–540. IEEE, 2021.
- [PS5] Ziqi Guo, Tingwen Bao, Wenlong Wu, Chao Jin, and Jay Lee. Iai devops: A systematic framework for prognostic model lifecycle management. In *2019 Prognostics and System Health Management Conference (PHM-Qingdao)*, pages 1–6. IEEE, 2019.
- [PS6] Rupesh Raj Karn, Prabhakar Kudva, and Ibrahim Abe M Elfadel. Dynamic autoselection and autotuning of machine learning models for cloud network analytics. *IEEE Transactions on Parallel and Distributed Systems*, 30(5):1052–1064, 2018.
- [PS7] Lucas Cardoso Silva, Fernando Rezende Zagatti, Bruno Silva Sette, Lucas Nildaimon dos Santos Silva, Daniel Lucrédio, Diego Furtado Silva, and Helena de Medeiros Caseli. Benchmarking machine learning solutions in production. In *2020 19th IEEE International Conference on Machine Learning and Applications (ICMLA)*, pages 626–633. IEEE, 2020.
- [PS8] Tamás Karácsony, Anna Mira Loesch-Biffar, Christian Vollmar, Soheyl Noachtar, and João Paulo Silva Cunha. Deepepil: Towards an epileptologist-friendly ai enabled seizure classification cloud system based on deep learning analysis of 3d videos. In *2021 IEEE EMBS International Conference on Biomedical and Health Informatics (BHI)*, pages 1–5. IEEE, 2021.
- [PS9] Chandrasekar Vuppapalapati, Anitha Ilapakurti, Karthik Chillara, Sharat Kedari, and Vanaja Mamidi. Automating tiny ml intelligent sensors devops using microsoft azure. In *2020 IEEE International Conference on Big Data (Big Data)*, pages 2375–2384. IEEE, 2020.
- [PS10] Tuomas Granlund, Aleksi Kopponen, Vlad Stirbu, Lalli Myllyaho, and Tommi Mikkonen. Mlops challenges in multi-organization setup: Experiences from two real-world cases. In *2021 IEEE/ACM 1st Workshop on AI Engineering-Software Engineering for AI (WAIN)*, pages 82–88. IEEE, 2021.
- [PS11] Lucy Ellen Lwakatare, Ivica Crnkovic, and Jan Bosch. Devops for ai—challenges in development of ai-enabled applications. In *2020 International Conference on Software, Telecommunications and Computer Networks (SoftCOM)*, pages 1–6. IEEE, 2020.
- [PS12] Benjamin Benni, Mireille Blay-Fornarino, Sébastien Mosser, Frédéric Précisio, and Günther Jungbluth. When devops meets meta-learning: A portfolio to rule them all. In *2019 ACM/IEEE 22nd International Conference on Model Driven Engineering Languages and Systems Companion (MODELS-C)*, pages 605–612. IEEE, 2019.
- [PS13] Ilias Gerostathopoulos, Stefan Kugele, Christoph Segler, Tomas Bures, and Alois Knoll. Automated trainability evaluation for smart software functions. In *2019 34th IEEE/ACM International Conference on Automated Software Engineering (ASE)*, pages 998–1001. IEEE, 2019.
- [PS14] Álvaro López García, Jesus Marco De Lucas, Marica Antonacci, Wolfgang Zu Castell, Mario David, Marcus Hardt, Lara Lloret Iglesias, Germán Moltó, Marcin Plociennik, Viet Tran, et al. A cloud-based framework for machine learning workloads and applications. *IEEE access*, 8:18681–18692, 2020.
- [PS15] Amine Barrak, Ellis E Eghan, and Bram Adams. On the co-evolution of ml pipelines and source code-empirical study of dvc projects. In *2021 IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER)*, pages 422–433. IEEE, 2021.
- [PS16] Fuyuki Ishikawa and Yutaka Matsuno. Evidence-driven requirements engineering for uncertainty of machine learning-based systems. In *2020 IEEE 28th International Requirements Engineering Conference (RE)*, pages 346–351. IEEE, 2020.
- [PS17] Markus Borg, Ronald Jabangwe, Simon Åberg, Arvid Ekblom, Ludwig Hedlund, and August Lidfeldt. Test automation with grad-cam heatmaps-a future pipe segment in mlops for vision ai? In *2021 IEEE International Conference on Software Testing, Verification and Validation Workshops (ICSTW)*, pages 175–181. IEEE, 2021.
- [PS18] Eero Kauhanen, Jukka K Nurminen, Tommi Mikkonen, and Matvei Pashkovskiy. Regression test selection tool for python in continuous integration process. In *2021 IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER)*, pages 618–621. IEEE, 2021.
- [PS19] Bojan Karlaš, Matteo Interlandi, Cedric Renggli, Wentao Wu, Ce Zhang, Deepak Mukunthu Iyappan Babu, Jordan Edwards, Chris Lauren, Andy Xu, and Markus Weimer. Building continuous integration services for machine learning. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 2407–2415, 2020.
- [PS20] Andrew Chen, Andy Chow, Aaron Davidson, Arjun DCunha, Ali Ghodsi, Sue Ann Hong, Andy Konwinski, Clemens Mewald, Siddharth Murching, Tomas Nykodym, et al. Developments in mlflow: A system to accelerate the machine learning lifecycle. In *Proceedings of the fourth international workshop on data management for end-to-end machine learning*, pages 1–4, 2020.
- [PS21] Cedric Renggli, Frances Ann Hubis, Bojan Karlaš, Kevin Schawinski, Wentao Wu, and Ce Zhang. Ease. ml/ci and ease. ml/meter in action: Towards data management for statistical generalization. *Proceedings of the VLDB Endowment*, 12(12):1962–1965, 2019.
- [PS22] Huaizheng Zhang, Yuanming Li, Yizheng Huang, Yonggang Wen, Jianxiong Yin, and Kyle Guan. Mlmodelci: An automatic cloud platform for efficient mlaas. In *Proceedings of the 28th ACM International Conference on Multimedia*, pages 4453–4456, 2020.
- [PS23] Maurits van der Goes. Scaling enterprise recommender systems for decentralization. In *Fifteenth ACM Conference on Recommender Systems*, pages 592–594, 2021.
- [PS24] Philipp Ruf, Manav Madan, Christoph Reich, and Djaffar Ould-Abdeslam. Demystifying mlops and presenting a recipe for the selection of open-source tools. *Applied Sciences*, 11(19):8861, 2021.

- [PS25] Emmanuel Raj, David Buffoni, Magnus Westerlund, and Kimmo Ahola. Edge mlops: An automation framework for aiot applications. In *2021 IEEE International Conference on Cloud Engineering (IC2E)*, pages 191–200. IEEE, 2021.
- [PS26] Silverio Martínez-Fernández, Xavier Franch, Andreas Jedlitschka, Marc Oriol, and Adam Trendowicz. Developing and operating artificial intelligence models in trustworthy autonomous systems. In *International Conference on Research Challenges in Information Science*, pages 221–229. Springer, 2021.
- [PS27] Romeo Kienzler and Ivan Nesic. Claimed, a visual and scalable component library for trusted ai. *arXiv preprint arXiv:2103.03281*, 2021.
- [PS28] Alexandra Posoldova. Machine learning pipelines: From research to production. *IEEE Potentials*, 39(6):38–42, 2020.
- [PS29] Yan Liu, Zhijing Ling, Boyu Huo, Boqian Wang, Tianen Chen, and Esma Mouine. Building a platform for machine learning operations from open source frameworks. *IFAC-PapersOnLine*, 53(5):704–709, 2020.
- [PS30] Mirella Sangiovanni, Gerard Schouten, and Willem-Jan van den Heuvel. An iot beehive network for monitoring urban biodiversity: vision, method, and architecture. In *Symposium and Summer School on Service-Oriented Computing*, pages 33–42. Springer, 2020.
- [PS31] Amitabha Banerjee, Chien-Chia Chen, Chien-Chun Hung, Xiaobo Huang, Yifan Wang, and Razvan Chevesaran. Challenges and experiences with mlops for performance diagnostics in hybrid-cloud enterprise software deployments. In *2020 USENIX Conference on Operational Machine Learning (OpML 20)*, 2020.
- [PS32] Willem-Jan van den Heuvel and Damian A Tamburri. Model-driven ml-ops for intelligent enterprise applications: vision, approaches and challenges. In *International Symposium on Business Modeling and Software Design*, pages 169–181. Springer, 2020.
- [PS33] Waldemar Hummer, Vinod Muthusamy, Thomas Rausch, Parijat Dube, Kaoutar El Maghraoui, Anupama Murthi, and Punleuk Oum. Modelops: Cloud-based lifecycle management for reliable and trusted ai. In *2019 IEEE International Conference on Cloud Engineering (IC2E)*, pages 113–120. IEEE, 2019.
- [PS34] Junsung Lim, Hoejoo Lee, Youngmin Won, and Hunje Yeon. Mlops lifecycle scheme for vision-based inspection process in manufacturing. In *2019 USENIX Conference on Operational Machine Learning (OpML 19)*, pages 9–11, 2019.
- [PS35] Behrouz Derakhshan, Alireza Rezaei Mahdiraji, Tilmann Rabl, and Volker Markl. Continuous deployment of machine learning pipelines. In *EDBT*, pages 397–408, 2019.
- [PS36] Stuart Jackson, Maha Yaqub, and Cheng-Xi Li. The agile deployment of machine learning models in healthcare. *Frontiers in big Data*, page 7, 2019.
- [PS37] Ola Spjuth, Jens Frid, and Andreas Hellander. The machine learning life cycle and the cloud: implications for drug discovery. *Expert opinion on drug discovery*, 16(9):1071–1079, 2021.
- [PS38] GeumSeong Yoon, Jungsu Han, Seunghyung Lee, and JongWon Kim. Devops portal design for smartx ai cluster employing cloud-native machine learning workflows. In *Advances in Internet, Data and Web Technologies: The 8th International Conference on Emerging Internet, Data and Web Technologies (EIDWT-2020)*, pages 532–539. Springer, 2020.
- [PS39] Krzysztof Czarnecki. Software engineering for automated vehicles: Addressing the needs of cars that run on software and data. In *2019 IEEE/ACM 41st International Conference on Software Engineering: Companion Proceedings (ICSE-Companion)*, pages 6–8. IEEE, 2019.
- [PS40] Jacopo Tagliabue. You do not need a bigger boat: Recommendations at reasonable scale in a (mostly) serverless and open stack. In *Proceedings of the 15th ACM Conference on Recommender Systems*, pages 598–600, 2021.
- [PS41] Philipp Kohl, Oliver Schmidts, Lars Klöser, Henri Werth, Bodo Kraft, and Albert Zündorf. Stamp 4 nlp—an agile framework for rapid quality-driven nlp applications development. In *International Conference on the Quality of Information and Communications Technology*, pages 156–166. Springer, 2021.
- [PS42] Aiswarya Raj Munappy, David Issa Mattos, Jan Bosch, Helena Holmström Olsson, and Anas Dakkak. From ad-hoc data analytics to dataops. In *Proceedings of the International Conference on Software and System Processes*, pages 165–174, 2020.
- [PS43] Haoyu Cai, Chao Wang, and Xuehai Zhou. Deployment and verification of machine learning tool-chain based on kubernetes distributed clusters: This paper is submitted for possible publication in the special issue on high performance distributed computing. *CCF Transactions on High Performance Computing*, 3(2):157–170, 2021.
- [PS44] Chaoyu Wu, E Haihong, and Meina Song. An automatic artificial intelligence training platform based on kubernetes. In *Proceedings of the 2020 2nd International Conference on Big Data Engineering and Technology*, pages 58–62, 2020.
- [PS45] Mlops: Methods and tools of devops for machine learning. Jul 2020. Online; accessed 1. Apr. 2022. <https://www.altexsoft.com/blog/mlops-methods-tools/>.
- [PS46] Youngjune Lee, Oh Joon Kwon, Haeju Lee, Joonyoung Kim, Kangwook Lee, and Kee-Eung Kim. Augment & value: A data enhancement pipeline for data-centric ai. *arXiv preprint arXiv:2112.03837*, 2021.
- [PS47] Tuomas Granlund, Vlad Stirbu, and Tommi Mikkonen. Towards regulatory-compliant mlops: Oravizio’s journey from a machine learning experiment to a deployed certified medical product. *SN computer Science*, 2(5):342, 2021.
- [PS48] Sindhu Ghanta, Sriram Subramanian, Lior Khermosh, Swaminathan Sundararaman, Harshil Shah, Yakov Goldberg, Drew Roselli, and Nisha Talagala. ML health monitor: taking the pulse of machine learning algorithms in production. In *Applications of Machine Learning*. International Society for Optics and Photonics, 2019.
- [PS49] Grigori Fursin, Herve Guillou, and Nicolas Essayan. Codereef: an open platform for portable mlops, reusable automation actions and reproducible benchmarking. *The Workshop on MLOps Systems at MLSys’20*, 2020.
- [PS50] Denis Baylor, Kevin Haas, Konstantinos Katsiapis, Sammy Leong, Rose Liu, Clemens Menwald, Hui Miao, Neoklis Polyzotis, Mitchell Trott, and Martin Zinkevich. Continuous training for production ml in the tensorflow extended (tfx) platform. In *2019 USENIX Conference on Operational*

*Machine Learning (OpML 19)*, pages 51–53, 2019.

- [PS51] Rama Akkiraju, Vibha Sinha, Anbang Xu, Jalal Mahmud, Pritam Gundecha, Zhe Liu, Xiaotong Liu, and John Schumacher. Characterizing machine learning processes: A maturity framework. In *International Conference on Business Process Management*, pages 17–31. Springer, 2020.
- [PS52] R Ciucu, FC Adochiei, Ioana-Raluca Adochiei, F Argatu, GC Seritan, B Enache, S Grigorescu, and Violeta Vasilica Argatu. Innovative devops for artificial intelligence. *The Scientific Bulletin of Electrical Engineering Faculty*, 19(1):58–63, 2019.
- [PS53] Emmanuel Raj, David Buffoni, Magnus Westerlund, and Kimmo Ahola. Edge mlops: An automation framework for aiot applications. In *IEEE International Conference on Cloud Engineering (IC2E)*, pages 191–200. IEEE, 2021.
- [PS54] Leonel Aguilar, David Dao, Shaoduo Gan, Nezihe Merve Gurel, Nora Hollenstein, Jiawei Jiang, Bojan Karlas, Thomas Lemmin, Tian Li, Yang Li, et al. Ease. ml: A lifecycle management system for mldev and mlops. *Proc. of Innovative Data Systems Research*, 2021.
- [PS55] Rohith Sothilingam, Eric Yu, and Arik Senderovich. Towards higher maturity for machine learning: A conceptual modelling approach. *The iJournal: Student Journal of the Faculty of Information*, 5(1):80–97, 2019.
- [PS56] Emmanuel Raj, Magnus Westerlund, and Leonardo Espinosa-Leal. Reliable fleet analytics for edge iot solutions. *arXiv preprint arXiv:2101.04414*, 2021.
- [PS57] Willem-Jan van den Heuvel and Damian A Tamburri. Model-driven ml-ops for intelligent enterprise applications: vision, approaches and challenges. In *Business Modeling and Software Design: 10th International Symposium, BMSD 2020, Berlin, Germany, July 6-8, 2020, Proceedings 10*, pages 169–181. Springer, 2020.
- [PS58] Cedric Renggli, Bojan Karlas, Bolin Ding, Feng Liu, Kevin Schawinski, Wentao Wu, and Ce Zhang. Continuous integration of machine learning models with ease. ml/ci: Towards a rigorous yet practical treatment. *Proceedings of Machine Learning and Systems*, 1:322–333, 2019.
- [PS59] Cedric Renggli, Luka Rimanic, Nezihe Merve Gürel, Bojan Karlas, Wentao Wu, and Ce Zhang. A data quality-driven view of mlops. *IEEE Data Engineering Bulletin*, 2021.
- [PS60] Fuyuki Ishikawa and Yutaka Matsuno. Continuous argument engineering: Tackling uncertainty in machine learning based systems. In *Computer Safety, Reliability, and Security: SAFECOMP 2018 Workshops, ASSURE, DECSO, SASSUR, STRIVE, and WAISE, Västerås, Sweden, September 18, 2018, Proceedings 37*, pages 14–21. Springer, 2018.
- [PS61] Christian Weber, Pascal Hirmer, Peter Reimann, and Holger Schwarz. A new process model for the comprehensive management of machine learning models. In *ICEIS (1)*, pages 415–422, 2019.
- [PS62] Denis Baylor, Eric Breck, Heng-Tze Cheng, Noah Fiedel, Chuan Yu Foo, Zakaria Haque, Salem Haykal, Mustafa Inspir, Vihan Jain, Levent Koc, et al. Tfx: A tensorflow-based production-scale machine learning platform. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1387–1395, 2017.
- [PS63] Tim Kraska, Ameet Talwalkar, John C Duchi, Rean Griffith, Michael J Franklin, and Michael I Jordan. Mlbase: A distributed machine-learning system. In *Cidr*, volume 1, pages 2–1, 2013.
- [PS64] Micah J Smith, Jürgen Cito, Kelvin Lu, and Kalyan Veeramachaneni. Enabling collaborative data science development with the ballet framework. *Proceedings of the ACM on Human-Computer Interaction*, 5(CSCW2):1–39, 2021.
- [PS65] Can Kaymakci, Simon Wenninger, and Alexander Sauer. A holistic framework for ai systems in industrial applications. In *International Conference on Wirtschaftsinformatik*, pages 78–93. Springer, 2021.
- [PS66] Pulkit Agrawal, Rajat Arya, Aanchal Bindal, Sandeep Bhatia, Anupriya Gagneja, Joseph Godlewski, Yucheng Low, Timothy Muss, Mudit Manu Paliwal, Sethu Raman, et al. Data platform for machine learning. In *Proceedings of the 2019 international conference on management of data*, pages 1803–1816, 2019.
- [PS67] Sridhar Alla and Suman Kalyan Adari. *Beginning MLOps with MLFlow*. Springer, 2021.
- [PS68] *MLOps for IT Teams: How to Transform the Machine Learning Lifecycle*. DataRobot, 2021.
- [PS69] *Practical MLOps - How to get ready for production models*. Valohai, 2020.
- [PS70] Niko Vuokko, Jukka Remes, Alexander Finn, Harry Souris, Hossein Yousefi, Joanna Purosto, and Pauliina Alanen. *Reliable and scalable AI with MLOps*. 2021.
- [PS71] Zechu Li, Xiao-Yang Liu, Jiahao Zheng, Zhaoran Wang, Anwar Walid, and Jian Guo. Finrl-podracr: High performance and scalable deep reinforcement learning for quantitative finance. *arXiv preprint arXiv:2111.05188*, 2021.
- [PS72] Vlad Stirbu, Tuomas Granlund, Jere Helén, and Tommi Mikkonen. Extending soup to ml models when designing certified medical systems. In *2021 IEEE/ACM 3rd International Workshop on Software Engineering for Healthcare (SEH)*, pages 32–35. IEEE, 2021.
- [PS73] Markus Borg. Agility in software 2.0—notebook interfaces and mlops with buttresses and rebars. In *International Conference on Lean and Agile Software Development*, pages 3–16. Springer, 2022.
- [PS74] Frances Ann Hubis, Wentao Wu, and Ce Zhang. Quantitative overfitting management for human-in-the-loop ml application development with ease. ml/meter. *arXiv preprint arXiv:1906.00299*, 2019.
- [PS75] Yizheng Huang, Huaizheng Zhang, Yonggang Wen, Peng Sun, and Nguyen Binh Duong TA. Modelci-e: Enabling continual learning in deep learning serving systems. *arXiv preprint arXiv:2106.03122*, 2021.
- [PS76] Dennis Muiruri, Lucy Ellen Lwakatare, Jukka K Nurminen, and Tommi Mikkonen. Practices and infrastructures for machine learning systems: An interview study in finnish organizations. *Computer*, 55(6):18–29, 2022.
- [PS77] Nikhil Muralidhar, Sathappah Muthiah, Patrick Butler, Manish Jain, Yu Yu, Katy Burne, Weipeng Li, David Jones, Prakash Arunachalam, Hays' Skip' McCormick, et al. Using antipatterns to avoid mlops mistakes. *arXiv preprint arXiv:2107.00079*, 2021.

- [PS78] Georgios Symeonidis, Evangelos Nerantzis, Apostolos Kazakis, and George A Papakostas. Mlops-definitions, tools and challenges. In *2022 IEEE 12th Annual Computing and Communication Workshop and Conference (CCWC)*, pages 0453–0460. IEEE, 2022.
- [PS79] Razvan Ciobanu, Alexandru Purdila, Laurentiu Piciu, and Andrei Damian. Solis—the mlops journey from data acquisition to actionable insights. *arXiv preprint arXiv:2112.11925*, 2021.
- [PS80] Xiao-Yang Liu, Zechu Li, Zhuoran Yang, Jiahao Zheng, Zhaoran Wang, Anwar Walid, Jian Guo, and Michael I Jordan. Elegantrl-podracr: Scalable and elastic library for cloud-native deep reinforcement learning. *arXiv preprint arXiv:2112.05923*, 2021.
- [PS81] Peizheng Li, Jonathan Thomas, Xiaoyang Wang, Ahmed Khalil, Abdelrahim Ahmad, Rui Inacio, Shipra Kapoor, Arjun Parekh, Angela Doufexi, Arman Shojaeifard, et al. Rlops: Development life-cycle of reinforcement learning aided open ran. *arXiv preprint arXiv:2111.06978*, 2021.
- [PS82] Chulhong Min, Akhil Mathur, Utku Günay Acer, Alessandro Montanari, and Fahim Kawsar. Sensix++: Bringing mlops and multi-tenant model serving to sensory edge devices. *ACM Transactions on Embedded Computing Systems*, 22(6):1–27, 2023.
- [PS83] Grigori Fursin. The collective knowledge project: Making ml models more portable and reproducible with open apis, reusable best practices and mlops. *arXiv preprint arXiv:2006.07161*, 2020.
- [PS84] Fan Yun, Toshiki Shibahara, Yuichi Ohsita, Daiki Chiba, Mitsuki Akiyama, and Masayuki Murata. Understanding machine learning model updates based on changes in feature attributions. 2020.
- [PS85] Christoph Windheuser, Danilo Sato, Alex Sadovsky, David Cooperberg, Greg Willis, Josh Poduska, Jim Falgout, and Dylan Tong. Mlops: Continuous delivery for machine learning on aws. 2020. (White Paper).
- [PS86] Mittal Anant Agrawal, Sumeet. Mlops: 5 steps to operationalize machine learning models. 2020. (White Paper).
- [PS87] Microsoft Corporation. Mlops with azure machine learning. Technical report, 2021.
- [PS88] Rohit Panikkar, Tamim Saleh, Maxime Szybowski, and Rob Whiteman. Operationalizing machine learning in processes. *McKinsey&Company*, September, 2021.
- [PS89] Khalid Salama, Jarek Kazmierczak, and Donna Schut. Practitioners guide to mlops: A framework for continuous delivery and automation of machine learning. 2021. (White Paper).
- [PS90] Building versus buying an ml management platform, 2020. Algorithmia Inc, (White Paper).
- [PS91] The complete guide to machine learning operations, 2020. run.ai, (White Paper).
- [PS92] Stephen Watts. What is mlops? machine learning operations explained. (2020). Retrieved April, 2022 from <https://www.bmc.com/blogs/mlops-machine-learning-ops>.
- [PS93] Harshit Tyagi. What is mlops? machine learning operations explained. (2021). Retrieved April, 2022 from <https://www.freecodecamp.org/news/what-is-mlops-machine-learning-operations-explained/>.
- [PS94] Alessandro Gaggia. Mlops essentials: four pillars for machine learning operations on aws. (2021). Retrieved April, 2022 from <https://www.proud2becloud.com/mlops-essentials-four-pillars-for-machine-learning-operations-on-aws/>.
- [PS95] Erin Wolpert. An overview of machine learning operations - building mature machine learning production systems through mlops. (2021). Retrieved April, 2022 from <https://developer.squareup.com/blog/an-overview-of-machine-learning-operations/>.
- [PS96] Ville Tuulos Outerbounds. Mlops vs. devops: Why data makes it different. (2021). Retrieved April, 2022 from <https://venturebeat.com/2021/10/30/mlops-vs-devops-why-data-makes-it-different/>.
- [PS97] Navdeep Singh Gill. Mlops platform - productionizing machine learning models. (2023). Retrieved November, 2023 from <https://www.xenonstack.com/blog/mlops/>.
- [PS98] Mlops: Methods and tools of devops for machine learning. (2020). Retrieved April, 2022 from <https://www.altexsoft.com/blog/mlops-methods-tools/>.
- [PS99] Seldon Technologies Ltd. Mlops software: Open source vs enterprise. (2022). Retrieved November, 2023 from <https://www.seldon.io/mlops-software-comparison>.
- [PS100] Danilo Sato, Arif Wider, and Christoph Windheuser. Continuous delivery for machine learning. (2019). Retrieved November, 2023 from <https://martinfowler.com/articles/cd4ml.html>.
- [PS101] Dominick Rocco. The 4 pillars of mlops: How to deploy ml models to production. (May 2021). Retrieved April, 2022 from <https://www.phdata.io/blog/the-ultimate-mlops-guide-how-to-deploy-ml-models-to-production>.
- [PS102] Google Cloud. Mlops: Continuous delivery and automation pipelines in machine learning. (January 2020). Retrieved April, 2022 from <https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning>.
- [PS103] Beatriz MA Matsui and Denise H Goya. Mlops: A guide to its adoption in the context of responsible ai. In *Proceedings of the 1st Workshop on Software Engineering for Responsible AI*, pages 45–49, 2022.
- [PS104] Rohith Sothilingam, Vik Pant, and SK Eric. Using i\* to analyze collaboration challenges in mlops project teams. In *iStar*, pages 1–6, 2022.
- [PS105] Xinrui Zhang and Jason Jaskolka. Conceptualizing the secure machine learning operations (secmlops) paradigm. In *2022 IEEE 22nd International Conference on Software Quality, Reliability and Security (QRS)*, pages 127–138. IEEE, 2022.
- [PS106] David Ping, Stefan Natu, Qingwei Li, Saeed Aghabozori, Vadim Dabravolski, Simon Zamarin, Ibrahim Gabr, and Amir Imani. Build a secure enterprise machine learning platform on aws, 2021. Amazon Web Services. AWS Technical Guide.
- [PS107] Bruno Klein. Aws prescriptive guidance - planning for successful mlops, 2021. Amazon Web Services.
- [PS108] Natalie Heard Vin Sharma, Rob Ferguson. Mlops: Emerging trends in data, code, and infrastructure., 2022. Amazon Web Services.
- [PS109] Canonical Ubuntu. A guide to mlops (white paper), 2023.

- [PS110] Databricks. How to automate your machine learning pipeline. (2022). Retrieved November 2023, from <https://www.databricks.com/resources/ebook/automate-your-machine-learning-pipeline>.
- [PS111] Amazon Web Services. Machine learning lens - aws well-architected framework. (2023). Retrieved November 2023, from <https://docs.aws.amazon.com/wellarchitected/latest/machine-learning-lens/machine-learning-lens.html>.
- [PS112] Sergio Zavota, Ajish Palakadan, Salman Taherian, and Mandeep Sehmi. Machine learning lens - aws well-architected framework, 2023. Amazon Web Services.
- [PS113] Databricks. What is mlops? (2023). Retrieved November 2023, from <https://www.databricks.com/glossary/mlops>.
- [PS114] Joseph Bradley, Rafi Kurlansik, Matt Thomson, and Niall Turbitt. The big book of mlops, 2023. Databricks, (E-book).
- [PS115] IBM Data and AI Team. Mlops and the evolution of data science. (2023). Retrieved November, 2023 from <https://www.ibm.com/think/topics/mlops>.
- [PS116] Elie Neghawi, Zerui Wang, Jun Huang, and Yan Liu. Linking team-level and organization-level governance in machine learning operations through explainable ai and responsible ai connector. In *2023 IEEE 47th Annual Computers, Software, and Applications Conference (COMPSAC)*, pages 1223–1230. IEEE, 2023.
- [PS117] Janvi Prasad, Arushi Jain, and Ushus Elizabeth Zachariah. Comparative evaluation of machine learning development lifecycle tools. In *2022 International Conference on Recent Trends in Microelectronics, Automation, Computing and Communications Systems (ICMACC)*, pages 1–6. IEEE, 2022.
- [PS118] Shailesh Kumar Shivakumar. Web performance monitoring and infrastructure planning. In *Modern Web Performance Optimization*, pages 175–212. Springer, 2020.
- [PS119] Seol Roh, Ki-Moon Jeong, Hye-Young Cho, and Eui-Nam Huh. An efficient microservices architecture for mlops. In *2023 Fourteenth International Conference on Ubiquitous and Future Networks (ICUFN)*, pages 652–654. IEEE, 2023.
- [PS120] Qinghua Lu, Liming Zhu, Xiwei Xu, Jon Whittle, David Douglas, and Conrad Sanderson. Software engineering for responsible ai: An empirical study and operationalised patterns. In *Proceedings of the 44th International Conference on Software Engineering: Software Engineering in Practice*, pages 241–242, 2022.
- [PS121] Gorka Zárate, Raúl Miñón, Josu Díaz-de Arcaya, and Ana I Torre-Bastida. K2e: Building mlops environments for governing data and models catalogues while tracking versions. In *2022 IEEE 19th International Conference on Software Architecture Companion (ICSA-C)*, pages 206–209. IEEE, 2022.
- [PS122] Sam Leroux, Pieter Simoens, Meelis Lootus, Kartik Thakore, and Akshay Sharma. Tinymlops: Operational challenges for widespread edge ai adoption. In *2022 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*, pages 1003–1010. IEEE, 2022.
- [PS123] Mariam Barry, Albert Bifet, and Jean-Luc Billy. Streamai: Dealing with challenges of continual learning systems for serving ai in production. In *2023 IEEE/ACM 45th International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP)*, pages 134–137. IEEE, 2023.
- [PS124] Mariam Barry, Jacob Montiel, Albert Bifet, Sameer Wadkar, Nikolay Manchev, Max Halford, Raja Chiky, Saad EL Jaouhari, Katherine B Shakman, Joudi Al Fehaily, et al. Streammlops: Operationalizing online learning for big data streaming & real-time applications. In *2023 IEEE 39th International Conference on Data Engineering (ICDE)*, pages 3508–3521. IEEE, 2023.
- [PS125] Ayush Shridhar and Deepak Nadig. Heuristic-based resource allocation for cloud-native machine learning workloads. In *2022 IEEE International Conference on Advanced Networks and Telecommunications Systems (ANTS)*, pages 415–418. IEEE, 2022.
- [PS126] Philipp Ruf, Christoph Reich, and Djaffar Ould-Abdeslam. Aspects of module placement in machine learning operations for cyber physical systems. In *2022 11th Mediterranean Conference on Embedded Computing (MECO)*, pages 1–6. IEEE, 2022.
- [PS127] Ilias Syrigos, Nikolaos Angelopoulos, and Thanasis Korakis. Optimization of execution for machine learning applications in the computing continuum. In *2022 IEEE Conference on Standards for Communications and Networking (CSCN)*, pages 118–123. IEEE, 2022.
- [PS128] Alexandre Carqueja, Bruno Cabral, João Paulo Fernandes, and Nuno Lourenço. On the democratization of machine learning pipelines. In *2022 IEEE Symposium Series on Computational Intelligence (SSCI)*, pages 455–462. IEEE, 2022.
- [PS129] Fotis Psallidas, Megan Eileen Leszczynski, Mohammad Hossein Namaki, Avriela Floratou, Ashvin Agrawal, Konstantinos Karanasos, Subru Krishnan, Pavle Subotic, Markus Weimer, Yinghui Wu, et al. Demonstration of geyser: Provenance extraction and applications over data science scripts. In *Companion of the 2023 International Conference on Management of Data*, pages 123–126, 2023.
- [PS130] Dezhao Tu, Yeye He, Weiwei Cui, Song Ge, Haidong Zhang, Shi Han, Dongmei Zhang, and Surajit Chaudhuri. Auto-validate by-history: Auto-program data quality constraints to validate recurring data pipelines. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 4991–5003, 2023.
- [PS131] Alec Gunny, Dylan Rankin, Philip Harris, Erik Katsavounidis, Ethan Marx, Muhammed Saleem, Michael Coughlin, and William Benoit. A software ecosystem for deploying deep learning in gravitational wave physics. In *Proceedings of the 12th Workshop on AI and Scientific Computing at Scale using Flexible Computing Infrastructures*, pages 9–17, 2022.
- [PS132] Steve Harris, Tim Bonnici, Thomas Keen, Watjana Lilaonitkul, Mark J White, and Nel Swanepoel. Clinical deployment environments: Five pillars of translational machine learning for health. *Frontiers in Digital Health*, 4:939292, 2022.
- [PS133] Tim Raffin, Tobias Reichenstein, Jonas Werner, Alexander Kühl, and Jörg Franke. A reference architecture for the operationalization of machine learning models in manufacturing. *Procedia CIRP*, 115:130–135, 2022.
- [PS134] Rustem Dautov, Erik Johannes Husom, and Fotis Gonidis. Towards mlops in mobile development with a plug-in architecture for data analytics. In *2022 6th International Conference on Computer, Software and Modeling (ICCSM)*, pages 22–27. IEEE, 2022.

- [PS135] Fabian Stieler and Bernhard Bauer. Git workflow for active learning - a development methodology proposal for data-centric ai projects, 18th international conference on evaluation of novel approaches to software engineering (enase 2023). 04 2023.
- [PS136] Xiaoyu Zhang, Jorge Piazentin Ono, Huan Song, Liang Gou, Kwan-Liu Ma, and Liu Ren. Sliceteller: A data slice-driven approach for machine learning model validation. *IEEE Transactions on Visualization and Computer Graphics*, 29(1):842–852, 2022.
- [PS137] Ruibo Chen and Wenjun Wu. Parallelizing automatic model management system for aiops on microservice platforms. In *European Conference on Parallel Processing*, pages 376–387. Springer, 2021.
- [PS138] John Wickström, Magnus Westerlund, and Emmanuel Raj. Decentralizing machine learning operations using web3 for iot platforms. In *2022 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*, pages 238–245. IEEE, 2022.
- [PS139] Nikos Psaromanolakis, Vasileios Theodorou, Dimitris Laskaratos, Ioannis Kalogeropoulos, Maria-Eleftheria Vlontzou, Eleni Zarogianni, and Georgios Samaras. Mlops meets edge computing: an edge platform with embedded intelligence towards 6g systems. In *2023 Joint European Conference on Networks and Communications & 6G Summit (EuCNC/6G Summit)*, pages 496–501. IEEE, 2023.
- [PS140] Alvaro Luis, Miguel A Patricio, Antonio Berlanga, and José M Molina. Seamless transition from machine learning on the cloud to industrial edge devices with thinger. io. *IEEE Internet of Things Journal*, 2023.
- [PS141] Guillermo Chinarro Álvarez, César Alejandro Achig Ramírez, Javier Andión, Juan C Dueñas, et al. Toward an integrated and supported machine learning process. In *2023 7th International Young Engineers Forum (YEF-ECE)*, pages 37–42. IEEE, 2023.
- [PS142] Sachchidanand Singh, Naveen Singh, and Vinay Singh. Comparative analysis of open standards for machine learning model deployments. In *ICT Systems and Sustainability: Proceedings of ICT4SD 2021, Volume 1*, pages 499–507. Springer, 2022.
- [PS143] Mattia Antonini, Miguel Pincheira, Massimo Vecchio, and Fabio Antonelli. Tiny-mlops: A framework for orchestrating ml applications at the far edge of iot systems. In *2022 IEEE international conference on evolving and adaptive intelligent systems (EAIS)*, pages 1–8. IEEE, 2022.
- [PS144] Hemadri Jayalath and Lakshmish Ramaswamy. Enhancing performance of operationalized machine learning models by analyzing user feedback. In *Proceedings of the 2022 4th International Conference on Image, Video and Signal Processing*, pages 197–203, 2022.
- [PS145] Laurent Boué, Pratap Kunireddy, and Pavle Subotić. Automatically resolving data source dependency hell in large scale data science projects. In *2023 IEEE/ACM 2nd International Conference on AI Engineering–Software Engineering for AI (CAIN)*, pages 1–6. IEEE, 2023.
- [PS146] DR Niranjana et al. Jenkins pipelines: A novel approach to machine learning operations (mlops). In *2022 International Conference on Edge Computing and Applications (ICECAA)*, pages 1292–1297. IEEE, 2022.
- [PS147] Inc. Red Hat. Mlops: Machine learning operations with red hat openshift. 2023. (White Paper).
- [PS148] Inc. Red Hat. Data, mlops, and iot for the next-generation insurance industry. Technical report, 2021.
- [PS149] Top 5 ways developers and data scientists can collaborate. (2022). Retrieved November, 2023 from <https://www.redhat.com/en/resources/top-5-ways-developers-collaborate-checklist>.
- [PS150] Nadia Nahar, Shurui Zhou, Grace Lewis, and Christian Kästner. Collaboration challenges in building ml-enabled systems: Communication, documentation, engineering, and process. In *Proceedings of the 44th international conference on software engineering*, pages 413–425, 2022.
- [PS151] Amazon AI. Amazon ai fairness and explainability. 2023. (White Paper).
- [PS152] Operationalizing machine learning. (2022). Retrieved April 2022 from <https://chicagodatascience.github.io/MLOps>.
- [PS153] Microsoft Azure. Machine learning operations (mlops). (2022). Retrieved April, 2022 from <https://azure.microsoft.com/en-au/products/machine-learning/mlops/>.
- [PS154] William Benton. Demystifying enterprise mlops. (2023). Retrieved March 2023 from <https://developer.nvidia.com/blog/demystifying-enterprise-mlops/>.
- [PS155] Souma Kanti Paul, Sadia Riaz, and Suchismita Das. A conceptual architecture for ai in supply chain risk management. In *TENCON 2022-2022 IEEE Region 10 Conference (TENCON)*, pages 1–5. IEEE, 2022.
- [PS156] Ruibo Chen, Yanjun Pu, Bowen Shi, and Wenjun Wu. An automatic model management system and its implementation for aiops on microservice platforms. *The Journal of Supercomputing*, 79(10):11410–11426, 2023.
- [PS157] Juan Pineda-Jaramillo and Francesco Viti. Mlops in freight rail operations. *Engineering Applications of Artificial Intelligence*, 123:106222, 2023.
- [PS158] Sebastian Schelter, Stefan Grafberger, Shubha Guha, Bojan Karlas, and Ce Zhang. Proactively screening machine learning pipelines with arguseyes. In *Companion of the 2023 International Conference on Management of Data*, pages 91–94, 2023.
- [PS159] Tobias Müller, Milena Zahn, and Florian Matthes. On the adoption of federated machine learning: Roles, activities and process life cycle. In *ICEIS (1)*, pages 525–531, 2023.
- [PS160] Rakshith Subramanya, Seppo Sierla, and Valeriy Vyatkin. From devops to mlops: Overview and application to electricity market forecasting. *Applied Sciences*, 12(19):9851, 2022.
- [PS161] Vlad Stirbu, Tuomas Granlund, and Tommi Mikkonen. Continuous design control for machine learning in certified medical systems. *Software Quality Journal*, 31(2):307–333, 2023.
- [PS162] Harsha Moraliyage, Dilantha Haputhanthri, Chamod Samarajeewa, Nishan Mills, Daswin De Silva, Milos Manic, and Andrew Jennings. Automated machine learning in critical energy infrastructure for net zero carbon emissions. In *2023 IEEE 32nd International Symposium on Industrial Electronics (ISIE)*, pages 1–7. IEEE, 2023.
- [PS163] Raúl Miñón, Josu Díaz-de Arcaya, Ana I Torre-Bastida, Gorka Zarate, and Aitor Moreno-Fernandez-de Leceta. Mlpacker: A unified software tool for packaging and deploying atomic and distributed analytic pipelines. In *2022 7th International Conference on Smart and Sustainable Technologies (SpliTech)*, pages 1–6. IEEE, 2022.

- [PS164] Marius Schlegel and Kai-Uwe Sattler. Mlflow2prov: extracting provenance from machine learning experiments. In *Proceedings of the Seventh Workshop on Data Management for End-to-End Machine Learning*, pages 1–4, 2023.
- [PS165] Dominik Kreuzberger, Julianne Forgo Kreuzberger, Mihai Criveti, Przemek Czuba, and Yvette Machowski. What is mlops? (2023). Retrieved November, 2023 from <https://ibm.github.io/MLOps/>.
- [PS166] David Nigenda, Zohar Karnin, Muhammad Bilal Zafar, Raghu Ramesha, Alan Tan, Michele Donini, and Krishnaram Kenthapadi. Amazon sagemaker model monitor: A system for real-time insights into deployed machine learning models. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 3671–3681, 2022.
- [PS167] Nathalie Rauschmayr, Sami Kama, Muhyun Kim, Miyoung Choi, and Krishnaram Kenthapadi. Profiling deep learning workloads at scale using amazon sagemaker. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 3801–3809, 2022.
- [PS168] Housseem Guissouma, Moritz Zink, and Eric Sax. Continuous safety assessment of updated supervised learning models in shadow mode. In *2023 IEEE 20th International Conference on Software Architecture Companion (ICSA-C)*, pages 301–308. IEEE, 2023.
- [PS169] Josu Diaz-de Arcaya, Ana I Torre-Bastida, Raúl Miñón, and Aitor Almeida. Orfeon: An aiops framework for the goal-driven operationalization of distributed analytical pipelines. *Future Generation Computer Systems*, 140:18–35, 2023.
- [PS170] Peini Liu, Gusseppe Bravo-Rocca, Jordi Guitart, Ajay Dholakia, David Ellison, and Miroslav Hodak. Scanflow-k8s: Agent-based framework for autonomic management and supervision of ml workflows in kubernetes clusters. In *2022 22nd IEEE International Symposium on Cluster, Cloud and Internet Computing (CCGrid)*, pages 376–385. IEEE, 2022.
- [PS171] Igor L Markov, Hanson Wang, Nitya S Kasturi, Shaun Singh, Mia R Garrard, Yin Huang, Sze Wai Celeste Yuen, Sarah Tran, Zehui Wang, Igor Glotov, et al. Looper: An end-to-end ml platform for product decisions. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 3513–3523, 2022.
- [PS172] Semo Yang, Jihwan Moon, Jinsoo Kim, Kwangkee Lee, and Kangyoon Lee. Flscalize: Federated learning lifecycle management platform. *IEEE Access*, 2023.
- [PS173] Hanqing Cao, Joshua Finer, Murad Meghani, Daniel C Nametz, Virginia Lorenzi, Lena Mamykina, Richard Meyers, Sarah C Rossetti, and Soojin Park. Machine learning model deployment using real-time physiological monitoring: Use case of detecting delayed cerebral ischemia. In *2022 IEEE Healthcare Innovations and Point of Care Technologies (HI-POCT)*, pages 42–45. IEEE, 2022.
- [PS174] Paul-Philipp Luley, Jan M Deriu, Peng Yan, Gerrit A Schatte, and Thilo Stadelmann. From concept to implementation: The data-centric development process for ai in industry. In *2023 10th IEEE Swiss Conference on Data Science (SDS)*, pages 73–76. IEEE, 2023.
- [PS175] Qi Cheng and Guodong Long. Federated learning operations (flops): Challenges, lifecycle and approaches. In *2022 International Conference on Technologies and Applications of Artificial Intelligence (TAAI)*, pages 12–17. IEEE, 2022.
- [PS176] Shir Chorev, Philip Tannor, Dan Ben Israel, Noam Bressler, Itay Gabbay, Nir Hutnik, Jonatan Liberman, Matan Perlmutter, Yurii Romanysyn, and Lior Rokach. Deepchecks: A library for testing and validating machine learning models and data. *Journal of Machine Learning Research*, 23(285):1–6, 2022.
- [PS177] Keita Sakuma, Ryuta Matsuno, and Yoshio Kameda. A method of identifying causes of prediction errors to accelerate mlops. In *2023 IEEE/ACM International Workshop on Deep Learning for Testing and Testing for Deep Learning (DeepTest)*, pages 9–16. IEEE, 2023.
- [PS178] Aquilas Tchanjou Njomou, Marios Fokaefs, Dimitry Fumtim Silatchom Kamga, and Bram Adams. On the challenges of migrating to machine learning life cycle management platforms. In *Proceedings of the 32nd Annual International Conference on Computer Science and Software Engineering*, pages 42–51, 2022.
- [PS179] Bo Wen, Yan Koyfman, Hongfei Tian, Boris Lublinsky, Raquel Norel, Carla Agurto, Dean Wampler, and Narciso Albarracin. Accelerating automation of digital health applications via cloud native approach. In *Proceedings of the Eighth International Workshop on Container Technologies and Container Clouds*, pages 1–6, 2022.
- [PS180] Gregor Cerar and Jernej Hribar. Machine learning operations model store: Optimizing model selection for ai as a service. In *2023 International Balkan Conference on Communications and Networking (BalkanCom)*, pages 1–5. IEEE, 2023.
- [PS181] Rustem Dautov, Erik Johannes Husom, Fotis Gonidis, Spyridon Papatzelos, and Nikolaos Malamas. Bridging the gap between java and python in mobile software development to enable mlops. In *2022 18th International Conference on Wireless and Mobile Computing, Networking and Communications (WiMob)*, pages 363–368. IEEE, 2022.
- [PS182] Qinghua Lu, Liming Zhu, Xiwei Xu, Jon Whittle, and Zhenchang Xing. Towards a roadmap on software engineering for responsible ai. In *Proceedings of the 1st International Conference on AI Engineering: Software Engineering for AI*, pages 101–112, 2022.
- [PS183] Rizal Broer Bahaweres, Aldi Zulfikar, Irman Hermadi, Arif Imam Suroso, and Yandra Arkeman. Docker and kubernetes pipeline for devops software defect prediction with mlops approach. In *2022 2nd International Seminar on Machine Learning, Optimization, and Data Science (ISMODE)*, pages 248–253. IEEE, 2022.
- [PS184] Bradley Eck, Duygu Kabakci-Zorlu, Yan Chen, France Savard, and Xiaowei Bao. A monitoring framework for deployed machine learning models with supply chain examples. In *2022 IEEE International Conference on Big Data (Big Data)*, pages 2231–2238. IEEE, 2022.
- [PS185] Rohan Gawhade, Lokesh Ramdev Bohara, Jesvin Mathew, and Poonam Bari. Computerized data-preprocessing to improve data quality. In *2022 Second International Conference on Power, Control and Computing Technologies (ICPC2T)*, pages 1–6. IEEE, 2022.
- [PS186] Claudia Patricia Ayala Martínez, Besim Bilalli, Cristina Gómez Seoane, and Silverio Juan Martínez Fernández. Dogo4ml: Development, operation and data governance for ml-based software systems. In *Joint Proceedings of RCIS 2022 Workshops and Research Projects Track: co-located with the 16th International Conference on Research Challenges in Information Science (RCIS 2022): Barcelona, Spain, May 17-20, 2022*. CEUR-WS. org, 2022.

- [PS187] Ashish Singh Parihar, Umesh Gupta, Utkarsh Srivastava, Vishal Yadav, and Vaibhav Kumar Trivedi. Automated machine learning deployment using open-source ci/cd tool. In *Proceedings of Data Analytics and Management: ICDAM 2022*, pages 209–222. Springer, 2023.
- [PS188] Raúl Miñón, Josu Diaz-de Arcaya, Ana I Torre-Bastida, and Philipp Hartlieb. Pangea: an mlops tool for automatically generating infrastructure and deploying analytic pipelines in edge, fog and cloud layers. *Sensors*, 22(12):4425, 2022.
- [PS189] Dhia Elhaq Rzig, Foyzul Hassan, Chetan Bansal, and Nachiappan Nagappan. Characterizing the usage of ci tools in ml projects. In *Proceedings of the 16th ACM/IEEE International Symposium on Empirical Software Engineering and Measurement*, pages 69–79, 2022.
- [PS190] Songzhu Mei, Cong Liu, Qinglin Wang, and Huayou Su. Model provenance management in mlops pipeline. In *Proceedings of the 2022 8th International Conference on Computing and Data Engineering*, pages 45–50, 2022.
- [PS191] Ayan Chatterjee, Bestoun S Ahmed, Erik Hallin, and Anton Engman. Quality assurance in mlops setting: An industrial perspective. *arXiv preprint arXiv:2211.12706*, 2022.
- [PS192] Joran Leest, Ilias Gerostathopoulos, and Claudia Raibulet. Evolvability of machine learning-based systems: An architectural design decision framework. In *2023 IEEE 20th International Conference on Software Architecture Companion (ICSA-C)*, pages 106–110. IEEE, 2023.
- [PS193] Antonio Guerriero, Roberto Pietrantuono, and Stefano Russo. Iterative assessment and improvement of dnn operational accuracy. In *2023 IEEE/ACM 45th International Conference on Software Engineering: New Ideas and Emerging Results (ICSE-NIER)*, pages 43–48. IEEE, 2023.
- [PS194] Georgios Samaras, Vasileios Theodorou, Dimitris Laskaratos, Nikolaos Psaromanolakis, Marinela Mertiri, and Alexandros Valantasis. Qmp: A cloud-native mlops automation platform for zero-touch service assurance in 5g systems. In *2022 IEEE International Mediterranean Conference on Communications and Networking (MeditCom)*, pages 86–89. IEEE, 2022.
- [PS195] Nils Baumann, Evgeny Kusmenko, Jonas Ritz, Bernhard Rumpe, and Moritz Benedikt Weber. Dynamic data management for continuous retraining. In *Proceedings of the 25th International Conference on Model Driven Engineering Languages and Systems: Companion Proceedings*, pages 359–366, 2022.
- [PS196] Jing Peng, Shiliang Zheng, Yutao Li, and Zhe Shuai. From source code to model service: A framework’s perspective. In *The International Conference on Natural Computation, Fuzzy Systems and Knowledge Discovery*, pages 1355–1362. Springer, 2022.
- [PS197] Markus Borg. Agility in software 2.0—notebook interfaces and mlops with buttresses and rebars. In *International Conference on Lean and Agile Software Development*, pages 3–16. Springer, 2022.
- [PS198] Eyad Kannout, Michał Grodzki, and Marek Grzegorowski. Considering various aspects of models’ quality in the ml pipeline-application in the logistics sector. In *2022 17th Conference on Computer Science and Intelligence Systems (FedCSIS)*, pages 403–412. IEEE, 2022.